

Detecting Citation Drift in AI-Generated Literature Summaries

Detecting citation drift in AI-generated literature summaries requires treating it as both a bibliographic validity problem and a claim-to-source alignment problem, because current systems often produce nonexistent references, corrupted metadata, and real citations that do not actually support the generated claims.

Abstract

Citation drift in AI-generated literature summaries refers to the divergence between a generated claim and the source it cites, including fabricated references, corrupted bibliographic metadata, overgeneralized interpretation, and misattribution of evidential support. This problem has become salient as large language models (LLMs) are increasingly used for literature review, related-work drafting, and evidence synthesis, while remaining prone to hallucination and factual inconsistency (Gao et al., 2023; Augenstein et al., 2023; Wang et al., 2025). Empirical studies show that AI systems can generate plausible-looking but nonexistent citations at substantial rates, with variation by model, domain, and prompt framing; even when cited works are real, metadata errors and unsupported attribution remain common (Walters & Wilder, 2023; Chelli et al., 2024; Oladokun et al., 2025; Erdem et al., 2024; Linardon et al., 2025). This paper synthesizes research on citation distortion, factual inconsistency in summarization, retrieval-augmented generation, and automated citation verification to define citation drift, organize its error modes, and review emerging detection methods. The evidence indicates that reliable detection requires multi-stage pipelines combining bibliographic resolution, evidence retrieval, claim decomposition, and entailment-style verification, while current benchmarks and metrics remain fragmented and domain-limited (Sarol et al., 2024; Yue et al., 2023; Yuan et al., 2026; Khajavi et al., 2026). The most important shift is from checking whether a citation exists to checking whether it faithfully supports a local claim. The biggest open question is how to build scalable, domain-robust, fine-grained evaluation systems for long-form scholarly synthesis.

Keywords

Citation drift; citation hallucination; large language models; literature summarization; retrieval-augmented generation; citation integrity; factual consistency; scholarly NLP.

Introduction

Scientific summarization aims to condense a growing research literature while preserving meaning, salience, and evidential support (Mohd et al., 2020; Mridha et al., 2021; Altmami & Menai, 2020). In practice, LLMs are now used for literature search, abstract screening, related-work drafting, and long-form review generation because manual synthesis is increasingly burdensome as publication volume expands (Wu et al., 2024; Orel et al., 2023; Tóth et al., 2024; Sen, 2025). Yet these systems remain vulnerable to hallucination, factual inconsistency, and opaque attribution, which makes them risky for scholarly synthesis (Wu et al., 2024; Augenstein et al., 2023; Wang et al., 2024).

“Citation drift” is a useful umbrella term for the ways a generated summary can separate from its evidence base. It includes at least three distinct failures: fabricated or non-resolving references, bibliographic mutation of real references, and semantic drift in which the cited source exists but does not support the generated claim as stated (Walters & Wilder, 2023; Mugaanyi et al., 2023; Khajavi et al., 2026; Peters & Chin-Yee, 2025). This framing connects recent AI citation-hallucination work to older meta-research on quotation error, citation distortion, and the propagation of unsupported claims through scientific networks (Greenberg, 2009; Cobb et al., 2023; Mogull, 2017).

Background and Related Work

Human scholarship already exhibits substantial citation distortion, which provides the conceptual baseline for AI-generated drift. In biomedical and psychological literatures, quotation or miscitation studies report nontrivial rates of unsupported, overgeneralized, or contradictory citing claims, with one review estimating 14.5% quotation error in medical research and a large psychology study finding about 19% of citing claims omit nuance or mischaracterize findings (Mogull, 2017; Cobb et al., 2023). Citation bias also systematically favors positive findings, which can amplify unsupported beliefs over time (Duyx et al., 2017).

Claim-specific citation-network studies show how bias, amplification, and invention can transform tentative or refuted claims into apparent facts through repeated citation alone (Greenberg, 2009). Similar dynamics appear in modern case studies, including persistent post-retraction citation and pandemic-era misrepresentation of null ophthalmology findings, showing that citation drift is not just a language-generation defect but a broader scholarly integrity problem (Schneider et al., 2020; Visioli et al., 2024).

A concise typology of citation drift in AI-generated summaries is useful here.

Drift Type	Operational Description	Representative Evidence
Fabrication	Citation points to no real scholarly work	High fabricated-reference rates in ChatGPT and Bard outputs (Walters & Wilder, 2023; Chelli et al., 2024)
Metadata mutation	Real target paper exists, but title, authors, year, DOI, or venue are corrupted	DOI hallucinations and partial matches are common, especially in humanities and specialized domains (Mugaanyi et al., 2023; Khajavi et al., 2026)
Attribution drift	Real citation does not support the generated claim	Generated citations often lack abstract-level substantiation, with near-zero citation accuracy in some long-form synthesis settings (Asai et al., 2026)
Interpretive drift	Summary broadens or alters study scope or conclusion	LLM summaries overgeneralize scientific conclusions more often than humans (Peters & Chin-Yee, 2025)
Network drift	Repetition diffuses unsupported claims through later writing	Citation cascades and post-retraction propagation sustain misinformation (Greenberg, 2009; Bar-Ilan & Halevi, 2018)

FIGURE 1 Types of citation drift in AI summaries

Main Technical Discussion

Detecting citation drift is harder than ordinary factuality checking because it is inherently multi-document and relational. Citation integrity requires linking a citing claim to the correct reference article, identifying the relevant evidence span within that article, and judging whether the article actually supports, contradicts, or fails to address the claim (Sarol et al., 2024). Citation-based scientific summarization and claim verification are therefore adjacent but incomplete antecedents: the former links citances to source text without assessing accuracy, and the latter often ignores the citing-document context and bibliographic layer (Sarol et al., 2024).

Current evidence suggests that bibliographic validity and semantic support must be separated analytically. Bibliographic checks ask whether a reference exists and whether its fields match authoritative records; semantic checks ask whether the cited source warrants the generated proposition (Yuan et al., 2026; Khajavi et al., 2026; Yue et al., 2023). This distinction matters because major hallucinations are often easier to detect than minor mutations or support mismatches (Khajavi et al., 2026).

A practical detection pipeline therefore has four stages.

Stage	Goal	Common Methods	Main Limitation
Citation resolution	Determine whether the citation maps to a real work	Database lookup, fuzzy matching, field-level comparison (Yuan et al., 2026)	Brittle to noisy or incomplete strings (Yuan et al., 2026)
Evidence retrieval	Find candidate passages supporting the claim	Scholarly retrieval, reranking, multi-source search (Sarol et al., 2024; Yuan et al., 2026)	Retrieval misses relevant support or surfaces distractors (Shao et al., 2023; Chen et al., 2023)
Claim attribution	Align local claim spans with evidence spans	Atomic claim decomposition, sentence-level citation generation (Xu et al., 2024; Xia et al., 2024)	Coarse paragraph-level citations mask local drift (Xu et al., 2024; Xia et al., 2024)
Support judgment	Decide whether evidence supports the claim	NLI, fact verification, structured LLM judging (Sarol et al., 2024; Yue et al., 2023; Wei et al., 2024)	Evaluators still miss subtle inconsistencies (Sarol et al., 2024; Chen et al., 2023)

FIGURE 2 Pipeline for detecting citation drift

LLM-generated summaries introduce a special form of interpretive drift even when references are real. In medical evidence summarization, models can produce factually inconsistent summaries, miss salient evidence, and become more error-prone with longer contexts (Tang et al., 2023). In scientific summarization more broadly, many models overgeneralize beyond source limits, and newer systems do not necessarily improve on this dimension (Peters & Chin-Yee, 2025). This means citation drift cannot be detected from reference strings alone.

Current Approaches and Developments

Empirical audits show large variation in citation reliability across models, disciplines, and prompts. GPT-3.5 fabricated 55% of citations in one multidisciplinary study, versus 18% for GPT-4, while substantive errors remained common even among real references (Walters & Wilder, 2023). In systematic-review prompting, hallucination rates ranged from 28.6% for GPT-4 to 91.4% for Bard, with low precision and recall against gold-standard review references (Chelli et al., 2024). Domain studies report similarly unstable performance in psychiatry, ophthalmology, librarianship, finance, mental health, and student writing (McGowan et al., 2023; Hua et al., 2023; Oladokun et al., 2025; Erdem et al., 2024; Linardon et al., 2025; Watson, 2024).

Recent systems increasingly use retrieval augmentation to reduce drift. RAG is motivated by the need to ground generation in external sources and improve verifiability (Gao et al., 2023; Huang & Huang, 2024; Wu et al., 2024). Benchmarks such as ALCE evaluate correctness and citation quality jointly, showing that even strong systems often lack complete citation support roughly half the time on difficult tasks (Gao et al., 2023). More recent grounded-generation methods report stronger attribution accuracy, including sentence-level approaches near 90% citation accuracy and self-reflective retrieval methods that improve citation accuracy in long-form generation (Xia et al., 2024; Asai et al., 2023; Ye et al., 2023).

Detection-specific systems are moving beyond exact matching.

System	Core Idea	Reported Strength
Citation Integrity corpus/model	Retrieve evidence sentences, then classify citation accuracy	Best model reached 0.59 micro-F1; subtle errors remain hard (Sarol et al., 2024)
CiteGuard	Retrieval-aware citation attribution alignment	68.1% on CiteME, near human 69.2% (Choi et al., 2025)
CiteAudit	Multi-agent bibliographic verification benchmark and pipeline	Stronger verification than commercial and open baselines (Yuan et al., 2026)
CiteCheck	Retrieve candidate work, then classify exact/minor/major drift	88.7 macro-F1 and 88.9% accuracy on physics benchmark (Khajavi et al., 2026)
String-feature classifier	Detect hallucinated references without external retrieval	AUC 0.876 cross-validation; useful as prescreening (Naser, 2026)

FIGURE 3 Recent systems for citation drift detection

Automated survey-generation systems are also beginning to bake in citation controls. LITERAS, OpenScholar, ScholarCopilot, AutoSurvey2, and DeepSurvey all combine retrieval with iterative or agentic verification and report improved citation fidelity relative to unguided generation (Gorenshtein et al., 2025; Asai et al., 2026; Wang et al., 2025; Wu et al., 2025; Yang et al., 2026). But retrieval alone does not solve attribution drift, because models still struggle with information integration, negative rejection, and false-information handling in RAG settings (Chen et al., 2023).

Challenges, Research Gaps, and Future Directions

The first challenge is conceptual: “citation accuracy” is often underdefined. Existing studies variously measure citation existence, metadata validity, DOI correctness, relevance to a topic, or support for a claim, which makes results hard to compare across systems (Chelli et al., 2024; Mugaanyi et al., 2023; Şahinuç et al., 2024). Benchmarks for scientific citation generation and verification remain fragmented across sentence-level, document-level, and long-form tasks (Funkquist et al., 2022; Xu et al., 2024).

The second challenge is evaluation. Standard summarization metrics do not capture factual consistency or citation support well, and automatic metrics often correlate weakly with expert judgments in high-stakes summarization (Tang et al., 2023; Davoodijam & Meybodi, 2024; Huang et al., 2021). LLM-based evaluators are promising, but current factuality judges remain imperfect, especially in domain-specific settings such as clinical summarization (Luo et al., 2024; Chen et al., 2023).

The third challenge is generalization across domains and prompt conditions. Citation fabrication rates vary by discipline, topic familiarity, and specificity, with poorer performance in underrepresented, long-tail, or specialized domains (Mugaanyi et al., 2023; Asai et al., 2026; Linardon et al., 2025). Lower-visibility or newer topics appear especially vulnerable to hallucination and metadata corruption (Erdem et al., 2024; Wu et al., 2024).

The fourth challenge is workflow integration. Manual validation remains necessary in evidence synthesis and AI-assisted review generation, and real-world adoption of automation is still limited outside screening tasks (Blaizot et al., 2022; Tóth et al., 2024). Because citations can function as data in reviews and bibliometric studies, hallucinated references can cross the line from poor scholarship into possible research misconduct when left unchecked (Resnik & Hosseini, 2026).

Priority future directions are clear.

Priority Gap	Why It Matters	Plausible Direction
Fine-grained claim-level benchmarks	Current metrics miss local attribution drift	Build datasets with atomic claims, source spans, and support labels (Xu et al., 2024; Chen et al., 2023)
Domain-robust verification	Performance drops on long-tail topics	Evaluate across humanities, biomedicine, and specialized subfields (Mugaanyi et al., 2023; Linardon et al., 2025)
End-to-end scholarly pipelines	Retrieval alone does not ensure faithful synthesis	Combine retrieval, decomposition, support judgment, and revision loops (Yang et al., 2026; Shao et al., 2023)
Publisher-side safeguards	Drift can enter preprints and publications	Automated screening, full-text deposit, and disclosure requirements (Walters & Wilder, 2023; Glynn, 2025; Linardon et al., 2025)

FIGURE 4 Research gaps in citation drift detection

Conclusion

Detecting citation drift in AI-generated literature summaries is not a single verification task but a layered problem spanning reference existence, metadata integrity, evidence retrieval, and claim-level support. The field's most important shift is from auditing whether a citation looks real to auditing whether each generated claim is faithfully grounded in an identifiable source. The biggest open question is how to scale that fine-grained verification across domains, long contexts, and multi-paper syntheses without making automated scholarly assistance too slow, brittle, or costly.

Complete References

Walters WH, Wilder EI. *Fabrication and errors in the bibliographic citations generated by ChatGPT*. **Scientific Reports**. 2023;13(1). DOI: 10.1038/s41598-023-41032-5 (Erdem et al., 2024).

Chelli M, Descamps J, Lavoué V, Trojani C, Azar M, Deckert M, Raynier J-L, Clowez G, Boileau P, Ruetsch-Chelli C. *Hallucination Rates and Reference Accuracy of ChatGPT and Bard for Systematic Reviews: Comparative Analysis*. **Journal of Medical Internet Research**. 2024;26:e53164. DOI: 10.2196/53164 (Erdem et al., 2024).

Gao T, Yen H, Yu J, Chen D. *Enabling Large Language Models to Generate Text with Citations*. 2023. Preprint/source as provided. Persistent identifier not listed in supplied record (Gao et al., 2023).

Mugaanyi J, Cai L, Cheng S-C, Lu C, Huang J. *Evaluation of Large Language Model Performance and Reliability for Citations and References in Scholarly Writing: Cross-Disciplinary Study*. **Journal of Medical Internet Research**. 2023/2024. DOI not fully provided in supplied record (Mugaanyi et al., 2023).

Sarol MJ, Ming S, Radhakrishna S, Schneider J, Kilicoglu H. *Assessing citation integrity in biomedical publications: corpus annotation and NLP models*. **Bioinformatics**. 2024. Repository: <https://github.com/ScienceNLP-Lab/Citation-Integrity/> (Sarol et al., 2024).

Greenberg SA. *How citation distortions create unfounded authority: analysis of a citation network*. **BMJ**. 2009. DOI not provided in supplied record (Greenberg, 2009).

Cobb CL, Crumly B, Montero-Zamora P, Schwartz SJ, Martínez CR. *The problem of miscitation in psychological science: Righting the ship*. **American Psychologist**. 2023. DOI not provided in supplied record (Cobb et al., 2023).

Mogull SA. *Accuracy of cited “facts” in medical research articles: A review of study methodology and recalculation of quotation error rate*. **PLoS ONE**. 2017. DOI not provided in supplied record (Mogull, 2017).

Peters U, Chin-Yee B. *Generalization bias in large language model summarization of scientific research*. **Royal Society Open Science**. 2025. DOI not provided in supplied record (Peters & Chin-Yee, 2025).

Tang L, Sun Z, Idnay B, Nestor J, Soroush A, Elias PA, Xu Z, Ding Y, Durrett G, Rousseau JF, Weng C, Peng Y. *Evaluating Large Language Models on Medical Evidence Summarization*. **medRxiv**. 2023. DOI not provided in supplied record (Tang et al., 2023).

Asai A, He J, Shao R, Shi W, Singh A, Chang J, Lo K, Soldaini L, Feldman S, D'Arcy M, Wadden D, et al. *Synthesizing scientific literature with retrieval-augmented language models*. **Nature**. 2026. DOI not provided in supplied record (Asai et al., 2026).

Yue X, Wang B, Zhang K, Chen Z, Su Y, Sun H. *Automatic Evaluation of Attribution by Large Language Models*. 2023. Preprint/source as provided (Yue et al., 2023).

Choi YM, Guo X, Fung YR, Wang Q. *CiteGuard: Faithful Citation Attribution for LLMs via Retrieval-Augmented Validation*. **arXiv**. 2025. Code: <https://github.com/KathCYM/CiteGuard> (Choi et al., 2025).

Yuan Z, Shi K, Zhang Z, Sun L, Chawla NV, Ye Y. *CiteAudit: You Cited It, But Did You Read It? A Benchmark for Verifying Scientific References in the LLM Era*. **arXiv**. 2026. Code: <https://github.com/shiiiikw/CiteAudit> (Yuan et al., 2026).

Khajavi K, Sadeghi S, Adhikari R, Tessier A. *CiteCheck: Retrieval-Grounded Detection of LLM Citation Hallucinations in Scientific Text*. **arXiv**. 2026. DOI/arXiv not provided in supplied record (Khajavi et al., 2026).

Ram O, Levine Y, Dalmedigos I, Muhlgay D, Shashua A, Leyton-Brown K, Shoham Y. *In-Context Retrieval-Augmented Language Models*. **Transactions of the Association for Computational Linguistics**. 2023. DOI not provided in supplied record (Ram et al., 2023).

Asai A, Wu Z, Wang Y, Sil A, Hajishirzi H. *Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection*. **arXiv**. 2023. arXiv/preprint as provided (Asai et al., 2023).

Xia S, Wang X, Liang J, Zhang Y, Zhou W, Deng J, Yu F, Xiao Y. *Ground Every Sentence: Improving Retrieval-Augmented LLMs with Interleaved Reference-Claim Generation*. 2024. Preprint/source as provided (Xia et al., 2024).

Huang Y-C, Feng X, Feng X, Qin B. *The Factual Inconsistency Problem in Abstractive Text Summarization: A Survey*. *arXiv*. 2021. arXiv/preprint as provided (Huang et al., 2021).

Resnik DB, Hosseini M. *Hallucinated citations produced by generative artificial intelligence may constitute research misconduct when citations function as data in scholarly papers*. *Accountability in Research*. 2026. DOI not provided in supplied record (Resnik & Hosseini, 2026).

These search results were found and analyzed using Consensus, an AI-powered search engine for research. Try it at <https://consensus.app>. © 2026 Consensus NLP, Inc. Personal, non-commercial use only; redistribution requires copyright holders' consent.

References

Altmami, N., & Menai, M. (2020). Automatic summarization of scientific articles: A survey. *Journal of King Saud University Computer and Information Sciences*, 34, 1011 - 1028. <https://doi.org/10.1016/j.jksuci.2020.04.020>

Asai, A., He, J., Shao, R., Shi, W., Singh, A., Chang, J., Lo, K., Soldaini, L., Feldman, S., D'Arcy, M., Wadden, D., Latzke, M., Sparks, J., Hwang, J. D., Kishore, V., Tian, M., Ji, P., Liu, S., Tong, H., . . . Hajishirzi, H. (2026). Synthesizing scientific literature with retrieval-augmented language models. *Nature*, 650, 857 - 863. <https://doi.org/10.1038/s41586-025-10072-4>

Asai, A., Wu, Z., Wang, Y., Sil, A., & Hajishirzi, H. (2023). Self-RAG: Learning to Retrieve, Generate, and Critique through Self-Reflection. *ArXiv, abs/2310.11511*. <https://doi.org/10.48550/arxiv.2310.11511>

Augenstein, I., Baldwin, T., Cha, M., Chakraborty, T., Ciampaglia, G. L., Corney, D., DiResta, R., Ferrara, E., Hale, S. A., Halevy, A., Hovy, E. H., Ji, H., Menczer, F., Míguez, R., Nakov, P., Scheufele, D. A., Sharma, S., & Zagni, G. (2023). Factuality challenges in the era of large language models and opportunities for fact-checking. *Nature Machine Intelligence*, 6, 852 - 863. <https://doi.org/10.1038/s42256-024-00881-z>

Bar-Ilan, J., & Halevi, G. (2018). Temporal characteristics of retracted articles. *Scientometrics*, 116, 1771-1783. <https://doi.org/10.1007/s11192-018-2802-y>

Blaizot, A., Veettil, S., Saidoung, P., Moreno-García, C. F., Wiratunga, N., Aceves-Martins, M., Lai, N., & Chaiyakunapruk, N. (2022). Using artificial intelligence methods for systematic review in health sciences: A systematic review. *Research Synthesis Methods*, 13, 353 - 362. <https://doi.org/10.1002/jrsm.1553>

Chelli, M., Descamps, J., Lavoué, V., Trojani, C., Azar, M., Deckert, M., Raynier, J., Clowez, G., Boileau, P., & Ruetsch-Chelli, C. (2024). Hallucination Rates and Reference Accuracy of ChatGPT and Bard for Systematic Reviews: Comparative Analysis. *Journal of Medical Internet Research*, 26. <https://doi.org/10.2196/53164>

Chen, J., Lin, H., Han, X., & Sun, L. (2023). Benchmarking Large Language Models in Retrieval-Augmented Generation. 17754-17762. <https://doi.org/10.48550/arxiv.2309.01431>

Chen, S., Zhao, Y., Zhang, J., Chern, E., Gao, S., Liu, P., & He, J. (2023). FELM: Benchmarking Factuality Evaluation of Large Language Models. *ArXiv, abs/2310.00741*. <https://doi.org/10.48550/arxiv.2310.00741>

Choi, Y. M., Guo, X., Fung, Y. R., & Wang, Q. (2025). CiteGuard: Faithful Citation Attribution for LLMs via Retrieval-Augmented Validation. *ArXiv, abs/2510.17853*. <https://doi.org/10.48550/arxiv.2510.17853>

Cobb, C. L., Crumly, B., Montero-Zamora, P., Schwartz, S. J., & Martínez, C. R. (2023). The problem of miscitation in psychological science: Righting the ship.. *The American psychologist*. <https://doi.org/10.1037/amp0001138>

- Davoodijam, E., & Meybodi, M. A. (2024). Evaluation metrics on text summarization: comprehensive survey. *Knowledge and Information Systems*, 66, 7717 - 7738. <https://doi.org/10.1007/s10115-024-02217-0>
- Duyx, B., Urlings, M. J. E., Swaen, G., Bouter, L., & Zeegers, M. (2017). Scientific citations favor positive results: a systematic review and meta-analysis.. *Journal of clinical epidemiology*, 88, 92-101. <https://doi.org/10.1016/j.jclinepi.2017.06.002>
- Erdem, O., Hassett, K., & Egriboyun, F. (2024). Hallucination in AI-generated financial literature reviews: evaluating bibliographic accuracy. *International Journal of Data Science and Analytics*, 20, 4501 - 4510. <https://doi.org/10.1007/s41060-025-00731-0>
- Funkquist, M., Kuznetsov, I., Hou, Y., & Gurevych, I. (2022). CiteBench: A benchmark for Scientific Citation Text Generation. 7337-7353. <https://doi.org/10.48550/arxiv.2212.09577>
- Gao, T., Yen, H., Yu, J., & Chen, D. (2023). Enabling Large Language Models to Generate Text with Citations. 6465-6488. <https://doi.org/10.48550/arxiv.2305.14627>
- Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., Guo, Q., Wang, M., & Wang, H. (2023). Retrieval-Augmented Generation for Large Language Models: A Survey. *ArXiv, abs/2312.10997*. <https://doi.org/10.48550/arxiv.2312.10997>
- Glynn, A. (2025). Guarding against artificial intelligence-hallucinated citations: the case for full-text reference deposit. *ArXiv, abs/2503.19848*. <https://doi.org/10.3897/ese.2025.e153973>
- Gorenshtein, A., Shihada, K., Sorka, M., Aran, D., & Shelly, S. (2025). LITERAS: Biomedical literature review and citation retrieval agents. *Computers in biology and medicine*, 192 Pt B, 110363. <https://doi.org/10.1016/j.combiomed.2025.110363>
- Greenberg, S. (2009). How citation distortions create unfounded authority: analysis of a citation network. *The BMJ*, 339. <https://doi.org/10.1136/bmj.b2680>
- Hua, H.-U., Kaakour, A.-H., Rachitskaya, A., Srivastava, S., Sharma, S., & Mammo, D. A. (2023). Evaluation and Comparison of Ophthalmic Scientific Abstracts and References by Current Artificial Intelligence Chatbots.. *JAMA ophthalmology*. <https://doi.org/10.1001/jamaophthalmol.2023.3119>
- Huang, Y., & Huang, J. (2024). A Survey on Retrieval-Augmented Text Generation for Large Language Models. *ACM Computing Surveys*, 58, 1 - 38. <https://doi.org/10.1145/3805774>
- Huang, Y.-C., Feng, X., Feng, X., & Qin, B. (2021). The Factual Inconsistency Problem in Abstractive Text Summarization: A Survey. *ArXiv, abs/2104.14839*. <https://doi.org/10.48550/arxiv.2104.14839>
- Khajavi, K., Sadeghi, S., Adhikari, R., & Tessier, A. (2026). CiteCheck: Retrieval-Grounded Detection of LLM Citation Hallucinations in Scientific Text. *ArXiv, abs/2605.27700*. <https://doi.org/10.48550/arxiv.2605.27700>
- Linardon, J., Jarman, H. K., McClure, Z., Anderson, C., Liu, C., & Messer, M. (2025). Influence of Topic Familiarity and Prompt Specificity on Citation Fabrication in Mental Health Research Using Large Language Models: Experimental Study. *JMIR Mental Health*, 12. <https://doi.org/10.2196/80371>
- Luo, Z., Xie, Q., & Ananiadou, S. (2024). Factual Consistency Evaluation of Summarisation in the Era of Large Language Models. *Expert Syst. Appl.*, 254, 124456. <https://doi.org/10.48550/arxiv.2402.13758>
- McGowan, A., Gui, Y., Dobbs, M., Shuster, S., Cotter, M., Selloni, A., Goodman, M., Srivastava, A., Cecchi, G., & Corcoran, C. (2023). ChatGPT and Bard exhibit spontaneous citation fabrication during psychiatry literature search. *Psychiatry research*, 326, 115334 - 115334. <https://doi.org/10.1016/j.psychres.2023.115334>

- Mogull, S. A. (2017). Accuracy of cited “facts” in medical research articles: A review of study methodology and recalculation of quotation error rate. *PLoS ONE*, 12. <https://doi.org/10.1371/journal.pone.0184727>
- Mohd, M., Jan, R., & Shah, M. (2020). Text document summarization using word embedding. *Expert Syst. Appl.*, 143. <https://doi.org/10.1016/j.eswa.2019.112958>
- Mridha, M. F., Lima, A. A., Nur, P. D. K., Das, S., Hasan, M., & Kabir, M. M. (2021). A Survey of Automatic Text Summarization: Progress, Process and Challenges. *IEEE Access*, PP, 1-1. <https://doi.org/10.1109/access.2021.3129786>
- Mugaanyi, J., Cai, L., Cheng, S.-C., Lu, C., & Huang, J. (2023). Evaluation of Large Language Model Performance and Reliability for Citations and References in Scholarly Writing: Cross-Disciplinary Study. *Journal of Medical Internet Research*, 26. <https://doi.org/10.2196/52935>
- Naser, M. (2026). How LLMs Cite and Why It Matters: A Cross-Model Audit of Reference Fabrication in AI-Assisted Academic Writing and Methods to Detect Phantom Citations. *ArXiv*, abs/2603.03299. <https://doi.org/10.48550/arxiv.2603.03299>
- Oladokun, B., Enakrire, R., Emmanuel, A., Ajani, Y., & Adetayo, A. J. (2025). Hallucitation in Scientific Writing: Exploring Evidence from ChatGPT Versions 3.5 and 4o in Responses to Selected Questions in Librarianship. *Journal of Web Librarianship*, 19, 62 - 92. <https://doi.org/10.1080/19322909.2025.2482093>
- Orel, E., Ciglenecki, I., Thiabaud, A., Temerev, A., Calmy, A., Keiser, O., & Merzouki, A. (2023). An Automated Literature Review Tool (LiteRev) for Streamlining and Accelerating Research Using Natural Language Processing and Machine Learning: Descriptive Performance Evaluation Study. *Journal of Medical Internet Research*, 25. <https://doi.org/10.2196/39736>
- Peters, U., & Chin-Yee, B. (2025). Generalization bias in large language model summarization of scientific research. *Royal Society Open Science*, 12. <https://doi.org/10.1098/rsos.241776>
- Ram, O., Levine, Y., Dalmedigos, I., Muhlgay, D., Shashua, A., Leyton-Brown, K., & Shoham, Y. (2023). In-Context Retrieval-Augmented Language Models. *Transactions of the Association for Computational Linguistics*, 11, 1316-1331. https://doi.org/10.1162/tacl_a_00605
- Resnik, D. B., & Hosseini, M. (2026). Hallucinated citations produced by generative artificial intelligence may constitute research misconduct when citations function as data in scholarly papers. *Accountability in research*, 2645390 - 2645390. <https://doi.org/10.1080/08989621.2026.2645390>
- Şahinuç, F., Kuznetsov, I., Hou, Y., & Gurevych, I. (2024). Systematic Task Exploration with LLMs: A Study in Citation Text Generation. *ArXiv*, abs/2407.04046. <https://doi.org/10.48550/arxiv.2407.04046>
- Sarol, M. J., Ming, S., Radhakrishna, S., Schneider, J., & Kilicoglu, H. (2024). Assessing citation integrity in biomedical publications: corpus annotation and NLP models. *Bioinformatics*, 40. <https://doi.org/10.1093/bioinformatics/btae420>
- Schneider, J., Ye, D., Ye, D., Hill, A. M., Whitehorn, A., & Whitehorn, A. (2020). Continued post-retraction citation of a fraudulent clinical trial report, 11 years after it was retracted for falsifying data. *Scientometrics*, 125, 2877 - 2913. <https://doi.org/10.1007/s11192-020-03631-1>
- Sen, C. K. (2025). Artificial Intelligence Tools in Biomedical Research: Part 1—Literature Search and Knowledge Mining. *Antioxidants & Redox Signaling*, 44, 1 - 10. <https://doi.org/10.1177/15230864251405885>
- Shao, Z., Gong, Y., Shen, Y., Huang, M., Duan, N., & Chen, W. (2023). Enhancing Retrieval-Augmented Large Language Models with Iterative Retrieval-Generation Synergy. *ArXiv*, abs/2305.15294. <https://doi.org/10.48550/arxiv.2305.15294>

- Tang, L., Sun, Z., Iday, B., Nestor, J., Soroush, A., Elias, P. A., Xu, Z., Ding, Y., Durrett, G., Rousseau, J. F., Weng, C., & Peng, Y. (2023). Evaluating Large Language Models on Medical Evidence Summarization. *medRxiv*, 6. <https://doi.org/10.1101/2023.04.22.23288967>
- Tóth, B., Berek, L., Gulácsi, L., Péntek, M., & Zrubka, Z. (2024). Automation of systematic reviews of biomedical literature: a scoping review of studies indexed in PubMed. *Systematic Reviews*, 13. <https://doi.org/10.1186/s13643-024-02592-3>
- Visioli, G., Pirraglia, M. P., Lambiase, A., & Gharbiya, M. (2024). Misleading Citations and Publication Bias in COVID-19 in Ophthalmology. *Journal of Evidence-Based Medicine*, 17, 703 - 704. <https://doi.org/10.1111/jebm.12664>
- Walters, W. H., & Wilder, E. I. (2023). Fabrication and errors in the bibliographic citations generated by ChatGPT. *Scientific Reports*, 13. <https://doi.org/10.1038/s41598-023-41032-5>
- Wang, C., Liu, X., Yue, Y., Guo, Q., Hu, X., Tang, X., Zhang, T., Jiayang, C., Yao, Y., Hu, X., Qi, Z., Gao, W., Wang, Y., Yang, L., Wang, J., Xie, X., Zhang, Z., & Zhang, Y. (2025). Survey on Factuality in Large Language Models. *ACM Computing Surveys*, 58, 1 - 37. <https://doi.org/10.1145/3742420>
- Wang, Y., Wang, M., Manzoor, M. A., Liu, F., Georgiev, G. N., Das, R. J., & Nakov, P. (2024). Factuality of Large Language Models: A Survey. 19519-19529. <https://doi.org/10.18653/v1/2024.emnlp-main.1088>
- Wang, Y., X., Nie, P., Zeng, H., Lyu, Z., Zhang, Y., Schneider, B., Lu, Y., Yue, X., & Chen, W. (2025). ScholarCopilot: Training Large Language Models for Academic Writing with Accurate Citations. *ArXiv*, abs/2504.00824. <https://doi.org/10.48550/arxiv.2504.00824>
- Watson, A. (2024). Hallucinated Citation Analysis: Delving into Student-Submitted AI-Generated Sources at the University of Mississippi. *The Serials Librarian*, 85, 172 - 180. <https://doi.org/10.1080/0361526x.2024.2433640>
- Wei, J., Yang, C., Song, X., Lu, Y., Hu, N., Huang, J., Tran, D., Peng, D., Liu, R., Huang, D., Du, C., & Le, Q. V. (2024). Long-form factuality in large language models. *ArXiv*, abs/2403.18802. <https://doi.org/10.48550/arxiv.2403.18802>
- Wu, S., X., Luo, D., Li, L., Shi, X., Chang, X., Lin, X., Luo, R., Pei, C., Zhao, Z., & Gong, J. (2024). Automated literature research and review-generation method based on large language models. *National Science Review*, 12. <https://doi.org/10.1093/nsr/nwaf169>
- Wu, S., Xiong, Y., Cui, Y., Wu, H., Chen, C., Yuan, Y., Huang, L., Liu, X., Kuo, T.-W., Guan, N., & Xue, C. (2024). Retrieval-augmented generation for natural language processing: a survey. *Artificial Intelligence Review*, 59. <https://doi.org/10.1007/s10462-026-11605-7>
- Wu, S., Liang, C., Bi, Z., Zhao, L., Wang, T., Song, J.-J., Zhang, Y., Chen, K., & Song, X. (2025). AutoSurvey2: Empowering Researchers with Next Level Automated Literature Surveys. *ArXiv*, abs/2510.26012. <https://doi.org/10.48550/arxiv.2510.26012>
- Xia, S., Wang, X., Liang, J., Zhang, Y., Zhou, W., Deng, J., Yu, F., & Xiao, Y. (2024). Ground Every Sentence: Improving Retrieval-Augmented LLMs with Interleaved Reference-Claim Generation. 969-988. <https://doi.org/10.48550/arxiv.2407.01796>
- Xu, Y., Gao, J., Yu, X., Bi, B., Shen, H., & Cheng, X. (2024). ALiCE: Evaluating Positional Fine-grained Citation Generation. *ArXiv*, abs/2406.13375. <https://doi.org/10.48550/arxiv.2406.13375>
- Yang, Z., D., Li, H., Wang, Z., Huang, T., Hu, Z., Wang, C., Zhang, Y., Wu, X., Yu, K., & Chen, L. (2026). DeepSurvey: Enhancing Analytical Depth and Citation Reliability in Automated Survey Generation. *ArXiv*, abs/2605.29522. <https://doi.org/10.48550/arxiv.2605.29522>

Ye, X., Sun, R., Arik, S. Ö., & Pfister, T. (2023). Effective Large Language Model Adaptation for Improved Grounding and Citation Generation. 6237-6251. <https://doi.org/10.48550/arxiv.2311.09533>

Yuan, Z., Shi, K., Zhang, Z., Sun, L., Chawla, N. V., & Ye, Y. (2026). CiteAudit: You Cited It, But Did You Read It? A Benchmark for Verifying Scientific References in the LLM Era. *ArXiv, abs/2602.23452*.
<https://doi.org/10.48550/arxiv.2602.23452>

Yue, X., Wang, B., Zhang, K., Chen, Z., Su, Y., & Sun, H. (2023). Automatic Evaluation of Attribution by Large Language Models. 4615-4635. <https://doi.org/10.48550/arxiv.2305.06311>